

DSI Summer Workshops Series

Peggy Lindner
Center for Advanced Computing & Data Science (CACDS)
Data Science Institute (DSI)
University of Houston
plindner@uh.edu

This Notebook is available at: <http://bitly.com/UHDSInotebook1> (<http://bitly.com/UHDSInotebook1>).

You can already download it! Use the "Save Link as" method.

Please make sure you have a copy of R up and running, as well as a Python 3 installation (ideally from Anacodna).

Goals for today

Understand basics of text analysis using R

(well enough so that you can Google your problems, find the answer, and implement it.)

More specifically

1. Up and running with R & IPython
2. Understand a basic exploratory data analysis workflow
3. Basics of R and Topic Modeling

Why R and not Python

It's good for data exploration!

Part 1: Getting yourself ready

First: Install software on your computer

- R [CRAN](https://www.anaconda.com/download/) (<https://www.anaconda.com/download/>)
- Python [Anaconda](https://www.anaconda.com/download/) (<https://www.anaconda.com/download/>)

Second: Prep your R environment

On a Mac open a terminal and start R

```
[plindner@peggys-mbp:~$ R

R version 3.5.0 (2018-04-23) -- "Joy in Playing"
Copyright (C) 2018 The R Foundation for Statistical Computing
Platform: x86_64-apple-darwin15.6.0 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

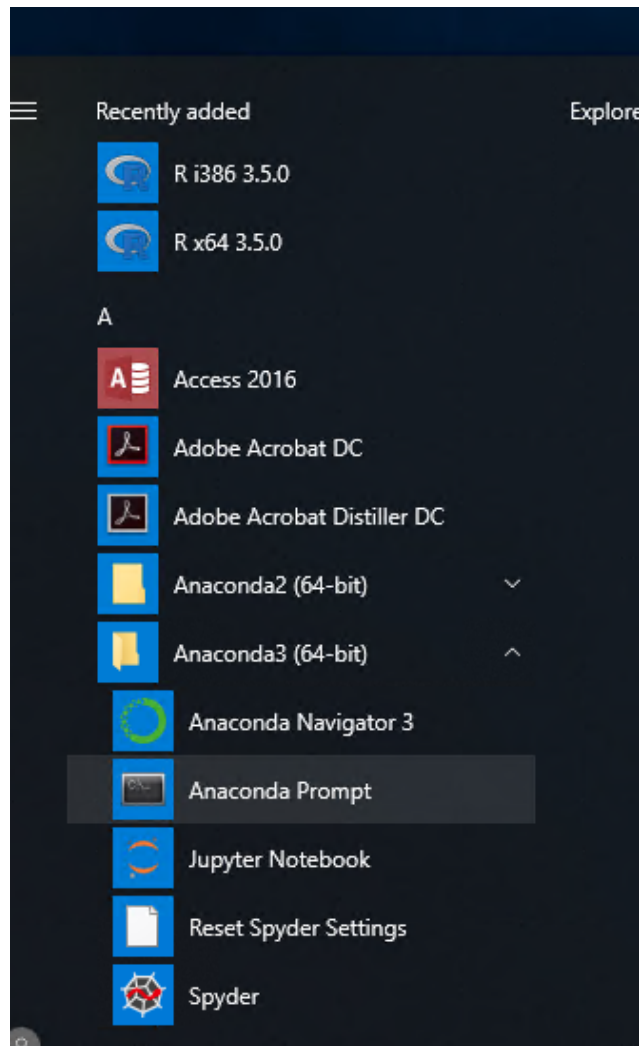
R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

[Previously saved workspace restored]

> █
```

On Windows: Open the Anaconda Command line and start R

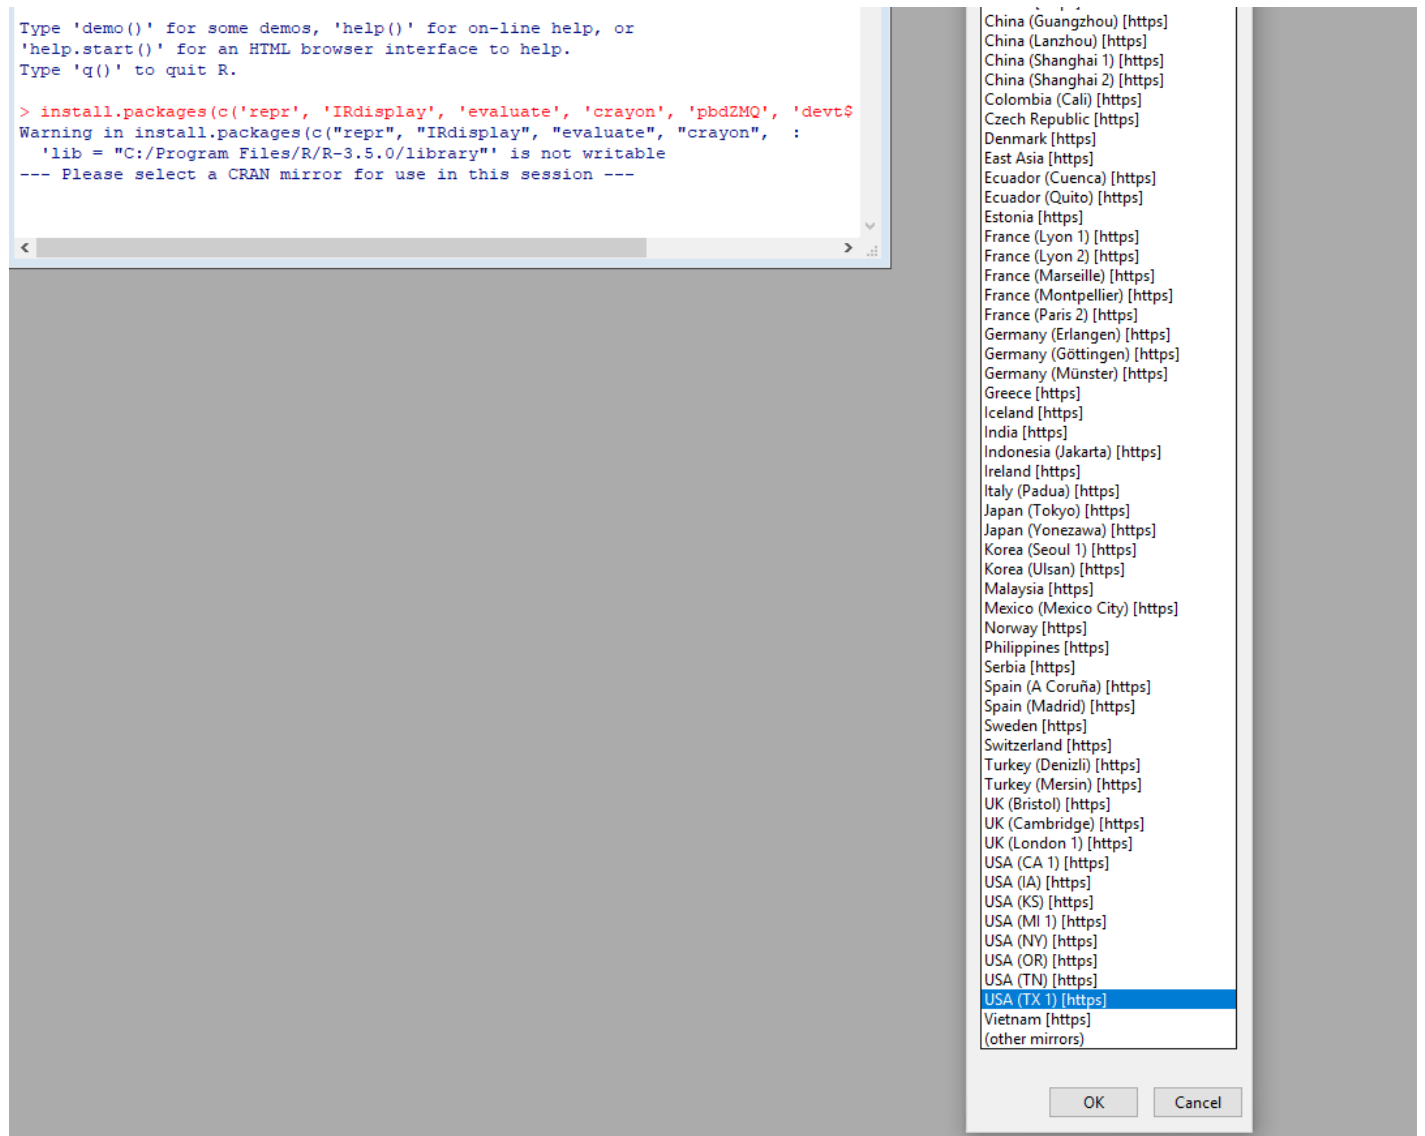


```
(C:\ProgramData\Anaconda3) C:\> cd C:\Program Files\R\R-3.5.0\bin\x64\  
(C:\ProgramData\Anaconda3) C:\Program Files\R\R-3.5.0\bin\x64>R  
  
R version 3.5.0 (2018-04-23) -- "Joy in Playing"  
Copyright (C) 2018 The R Foundation for Statistical Computing  
Platform: x86_64-w64-mingw32/x64 (64-bit)  
  
R is free software and comes with ABSOLUTELY NO WARRANTY.  
You are welcome to redistribute it under certain conditions.  
Type 'license()' or 'licence()' for distribution details.  
  
Natural language support but running in an English locale  
  
R is a collaborative project with many contributors.  
Type 'contributors()' for more information and  
'citation()' on how to cite R or R packages in publications.  
  
Type 'demo()' for some demos, 'help()' for on-line help, or  
'help.start()' for an HTML browser interface to help.  
Type 'q()' to quit R.  
  
> IRkernel::installspec()  
[InstallKernelSpec] Installed kernelspec ir in C:\Users\plindner\AppData\Roaming\jupyter\kernels\ir  
> q()
```


Now let's install some packages ...

```
> install.packages(c('readr', 'stringr', 'SnowballC', 'wordcloud', 'RColorBrewer'))
> install.packages(c('tm', 'ggplot2', 'topicmodels'))
> install.packages(c('repr', 'IRdisplay', 'evaluate', 'crayon', 'pbdZMQ', 'devtools', 'uuid', 'digest'))
> devtools::install_github('IRkernel/IRkernel')
```

When you see "Please select a CRAN mirror" , well select one.



... one last step - installing the Kernel

```
> IRkernel::installspec()
```

Now we can close the R environment (but leave your terminal and console open)

```
> quit()
```

Say "N" (no) when asked to save the workspace.

Jupyter Notebooks is what we will be going to use

We are now ready to start up our Jupyter Environment from the terminal or the console:

```
$ jupyter notebook --notebook-dir C:/Users/[your username]
```

or on a Mac

```
$ jupyter notebook --notebook-dir /Users/[your username]
```

And your browser should open at the address: <http://localhost:8888/tree> (<http://localhost:8888/tree>)



Open the downloaded notebook on your computer

 Quit Logout

Files Running Clusters

Select items to perform actions on them. Upload New ▾ 

0 ▾ / Downloads Name Last Modified ↑ File size

..	seconds ago	
 DSI Workshop Summer 2018a.ipynb	11 minutes ago	6.99 kB
 Simple_0.6.tar.gz	7 hours ago	223 kB
 Simple	7 hours ago	

Quick intro to Jupyter notebooks

Cells can be Markdown (like this one) or code

To start off with

Make sure you hit Shift-Enter or Ctrl-Enter when you are done.

In []: `2 + 2`

Part 2: The Exploratory Analysis Workflow

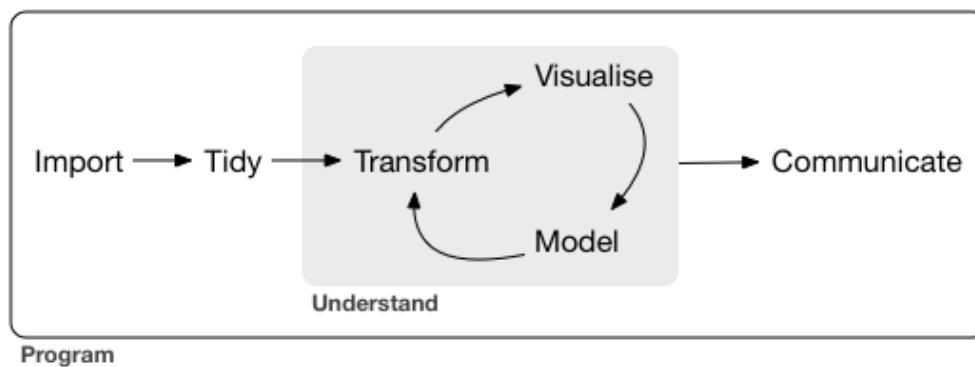


Image source: Hadley Wickham, R for Data Science

Our Example

Media Analysis of a bunch of articles downloaded from a database called "Factiva"

The screenshot shows the Factiva web interface. The top navigation bar includes 'Home', 'Search', 'Alerts', 'News Pages', and 'Companies/Markets'. The 'Search' bar is active, showing filters for 'DATE: In the last 3 months' and 'SOURCE: All Sources'. Below the search bar, there are tabs for 'Dow Jones (40)', 'All (471)', 'Publications', 'Web News (264)', 'Blogs (12)', 'Pictures (0)', and 'Multimedia (2)'. The 'Publications' tab is selected. On the left, there are filters for 'Sources' and 'Subjects'. The 'Sources' filter shows a list of sources like 'Dow Jones Newswires', 'Daily Mail', 'Financial Times', etc. The 'Subjects' filter shows 'Profit Warnings', 'Earnings', and 'Management Moves'. The main content area displays a list of search results. The first result is titled 'ALEX BRUMMER: Balfour Beatty tunnels a black hole with fifth profit warning since 2012'. The second result is titled 'Fitch: Tesco Warning Highlights Need for Rebate Disclosures'. The third result is also titled 'Fitch: Tesco Warning Highlights Need for Rebate Disclosures'. The fourth result is titled 'Fears Tesco profit warnings could lead to further construction cuts'. On the right, there is a 'Search Summary' table with the following data:

Search Summary	
Text	
Date	In the last 3 months
Source	All Sources
Author	All Authors
Company	Tesco PLC
Subject	Profit Warnings
Industry	All Industries
Region	All Regions
Language	English
Results Found	471
Timestamp	30 September 2014 8:44

Make sure you download the data source file: <http://bit.ly/UHDSldata1> (<http://bit.ly/UHDSldata1>) and store it in the Jupyter notebook directory (next to the *.ipynb file)

Frequently used R Packages in conjunction with text data

- [readr](https://cran.r-project.org/web/packages/readr/readr.pdf) (<https://cran.r-project.org/web/packages/readr/readr.pdf>) Import data

Data Analysis of text based material

- [stringr](https://cran.r-project.org/web/packages/stringr/vignettes/stringr.html) (<https://cran.r-project.org/web/packages/stringr/vignettes/stringr.html>) Clean up text
- [SnowballC](https://cran.r-project.org/web/packages/SnowballC/SnowballC.pdf) (<https://cran.r-project.org/web/packages/SnowballC/SnowballC.pdf>) Stemming of words
- [tm](https://cran.r-project.org/web/packages/tm/vignettes/tm.pdf) (<https://cran.r-project.org/web/packages/tm/vignettes/tm.pdf>) Text mining
- [Quanteda](https://quanteda.io/) (<https://quanteda.io/>) versatile text analysis tool
- [topicmodels](https://www.tidytextmining.com/topicmodeling.html) (<https://www.tidytextmining.com/topicmodeling.html>) Topic Modeling

Visualization

- [ggplot2](http://ggplot2.tidyverse.org/) (<http://ggplot2.tidyverse.org/>) Modern R visualizations
- [wordcloud](http://developer.marvel.com) (<http://developer.marvel.com>) Make some nice word clouds
- [RColorBrewer](https://dataset.readthedocs.org/en/latest/) (<https://dataset.readthedocs.org/en/latest/>) Get color into your visualizations


```
In [ ]: #load all required libraries
library(readr)
library(stringr)
library(SnowballC)
library(wordcloud)
library(RColorBrewer)
library(tm)
library(ggplot2)
library(topicmodels)
```

Data Import

```
In [ ]: # put the name of your csv file
inputfile <- "AJA_Factiva.txt"
# read the data
alldata <- read_file(inputfile)
# look at the dat
# what type is our data?
str(alldata)
```

Prepare data

```
In [ ]: # data wrangling - split the file in different articles
split.word <- "Document AJAZEN(.*)"

# split up into individual documents
list_alldataSplitted <- str_split(alldata, split.word)
# convert to vector and remove last element (which is a leftover)
alldataSplitted <- unlist(list_alldataSplitted)
alldataSplitted <- alldataSplitted[-length(alldataSplitted)]
str(alldataSplitted)
```

```
In [ ]: ### create corpus
article.corpus <- Corpus(VectorSource((alldataSplitted)))

article.corpus
```

```
In [ ]: #inspect a particular document
writeLines(as.character(article.corpus[[30]]))
```

```
In [ ]: #Check details (look at bunched up corpus to find anomalies)
inspect(article.corpus)
```

Data cleaning

```
In [ ]: #create the toSpace content transformer
toSpace <- content_transformer(function(x, pattern) { return (gsub(pattern, "
", x))})
#to remove potentially problematic symbols
article.corpus <- tm_map(article.corpus, toSpace, "-")
article.corpus <- tm_map(article.corpus, toSpace, ":")
article.corpus <- tm_map(article.corpus, toSpace, "'")
article.corpus <- tm_map(article.corpus, toSpace, "\"")
article.corpus <- tm_map(article.corpus, toSpace, " -")

#Good practice to check after each step.
writeLines(as.character(article.corpus[[30]]))
```

```
In [ ]: #Remove punctuation - replace punctuation marks with " "
article.corpus <- tm_map(article.corpus, removePunctuation)

#Good practice to check after each step.
writeLines(as.character(article.corpus[[30]]))
```

```
In [ ]: #Transform to lower case
article.corpus <- tm_map(article.corpus, content_transformer(tolower))

#Strip digits
article.corpus <- tm_map(article.corpus, removeNumbers)

#Remove stopwords from standard stopwords list
article.corpus <- tm_map(article.corpus, removeWords, stopwords("english"))

#inspect output
writeLines(as.character(article.corpus[[30]]))
```

Stopwords (<https://github.com/arc12/Text-Mining-Weak-Signals/wiki/Standard-set-of-english-stopwords>)

```
In [ ]: #define and eliminate all custom stopwords
myStopwords <- c("monday")
article.corpus <- tm_map(article.corpus, removeWords, myStopwords)

#Strip whitespace (cosmetic?)
article.corpus <- tm_map(article.corpus, stripWhitespace)

#inspect output
writeLines(as.character(article.corpus[[30]]))
```

Word Stemming (<http://www.omegahat.net/Rstem/stemming.pdf>)

```
In [ ]: #Stem document
article.corpus <- tm_map(article.corpus, stemDocument)

#inspect output
writeLines(as.character(article.corpus[[30]]))
```

Prepare for Analysis - create word counts

```

In [ ]: #Create document-term matrix
dtm <- DocumentTermMatrix(article.corpus)

dtm

In [ ]: #inspect segment of document term matrix
inspect(dtm[15:16,100:105])

In [ ]: #collapse matrix by summing over columns - this gets total counts (over all d
ocs) for each term
freq <- colSums(as.matrix(dtm))
#length should be total number of terms
length(freq)

In [ ]: #create sort order (descending)
ord <- order(freq,decreasing=TRUE)
#inspect most frequently occurring terms
freq[head(ord)]
#inspect least frequently occurring terms
freq[tail(ord)]

#List all terms in decreasing order of freq and write to disk
write.csv(freq[ord],"word_freq.csv")

In [ ]: #alterantive: remove very frequent and very rare words
dtmr <-DocumentTermMatrix(article.corpus, control=list(wordLengths=c(4, 20),
                bounds = list(global = c(3,27))))

dtmr

freqr <- colSums(as.matrix(dtmr))
#length should be total number of terms
length(freqr)

#create sort order (desc)
ordr <- order(freqr,decreasing=TRUE)
#inspect most frequently occurring terms
freqr[head(ordr)]
#inspect least frequently occurring terms
freqr[tail(ordr)]

In [ ]: #list most frequent terms. Lower bound specified as second argument
findFreqTerms(dtmr,lowfreq=60)

```

Now that we have the most frequently occurring terms in hand, we can check for correlations between some of these and other terms that occur in the corpus. In this context, correlation is a quantitative measure of the co-occurrence of words in multiple documents.

```

In [ ]: #correlations
findAssocs(dtmr,"turkish",0.5)
findAssocs(dtm,"children",0.5)

```

One needs to specify the DTM, the term of interest and the correlation limit. The latter is a number between 0 and 1 that serves as a lower bound for the strength of correlation between the search and result terms. For example, if the correlation limit is 1, findAssocs() will return only those words that always co-occur with the search term. A correlation limit of 0.5 will return terms that have a search term co-occurrence of at least 50% and so on.

Visualizations

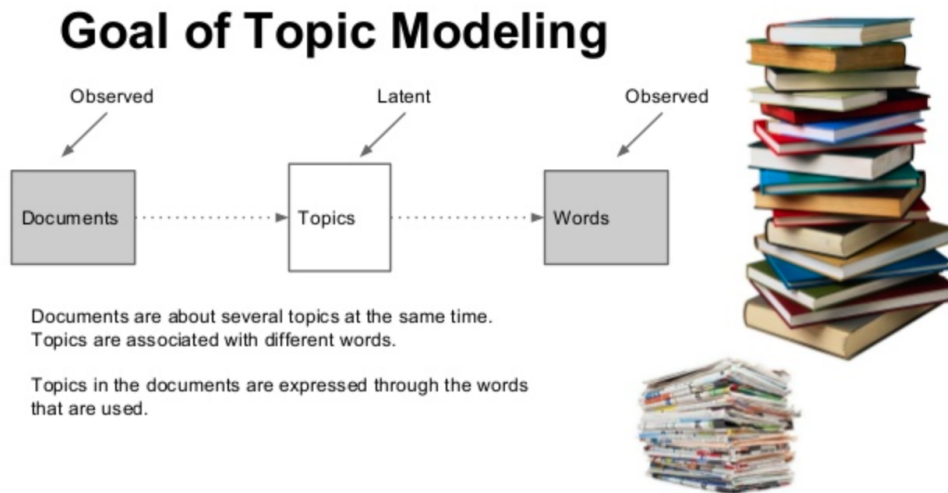
```
In [ ]: #Basic graphics
        #histogram
        wf=data.frame(term=names(freq),occurrences=freq)

        p <- ggplot(subset(wf, freq>100), aes(term, occurrences))
        p <- p + geom_bar(stat="identity")
        p <- p + theme(axis.text.x=element_text(angle=45, hjust=1))
        p
```

```
In [ ]: #wordcloud
        #setting the same seed each time ensures consistent look across clouds
        set.seed(42)
        #limit words by specifying min frequency
        wordcloud(names(freq),freq, min.freq=70)
        #...add color
        wordcloud(names(freq),freq,min.freq=70,colors=brewer.pal(6,"Dark2"))
```

What is Topic Modeling

- deals with the problem of automatically classifying sets of documents into themes



What is behind Topic Modeling?

```
In [ ]: #Topic modeling
        #Set parameters for Gibbs sampling
        burnin <- 4000
        iter <- 2000
        thin <- 500
        seed <-list(2003,5,63,100001,765)
        nstart <- 5
        best <- TRUE

        #Number of topics
        k <- 5

        #Run LDA using Gibbs sampling
        ldaOut <-LDA(dtm,k, method="Gibbs", control=list(nstart=nstart, seed = seed,
        best=best, burnin = burnin, iter = iter, thin=thin))
```

```
In [ ]: #have a look at the model and some output
        ldaOut
        topics(ldaOut)
        as.matrix(terms(ldaOut,6))

        #write out results
        #docs to topics
        ldaOut.topics <- as.matrix(topics(ldaOut))
        write.csv(ldaOut.topics,file=paste("LDAGibbs",k,"DocsToTopics.csv"))

        #top 6 terms in each topic
        ldaOut.terms <- as.matrix(terms(ldaOut,6))
        write.csv(ldaOut.terms,file=paste("LDAGibbs",k,"TopicsToTerms.csv"))
```